


INTERACTIVE VISUAL GYM / FIGURE PROOF

Environment showcase / sampled sequences | 5.5-inch manuscript width



Procedural tasks and sampled observations. Each row pairs a task description with selected user-view checkpoints. Top-down maps provide context, not execution steps. Ellipses omit actions. The phone row shows complete screens; hand and drag annotations illustrate the actual simulator gestures.

This proof places the figure at the manuscript text width, followed by a 10-point caption. Screenshot details do not replace readable vector labels. The publication PDF preserves text and geometry as vectors; only simulator observations are raster images.

INTERACTIVE VISUAL GYM / FIGURE PROOF

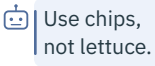
Interaction teaser / episode excerpts | 5.5-inch manuscript width

Simulated human error: wrong ingredient — lettuce instead of fried potato

Assistant A ...



Approach lettuce




Use chips, not lettuce.

| Oh, right.




Go to fryer → scoop chips



Now take it to the pass.

Go to pass → serve

Assistant C ...



I (silent)



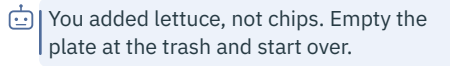
Scoop lettuce → go to pass

Serve → rejected

Assistant B Two user continuations



Scoop lettuce



You added lettuce, not chips. Empty the plate at the trash and start over.

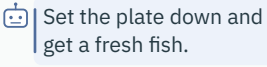
High trust
Welcomes step-by-step guidance

| Got it, emptying the plate.



Go to trash → empty plate

| What's next?




Set the plate down and get a fresh fish.



Put down plate → fetch fish ...

Low trust
Dislikes interruptions

| Why throw it out? It's just lettuce.



The order needs fish and chips, with no lettuce.

| Fine, I'll redo it.



Go to trash → empty plate



Get a fresh fish.

| I know the rest.

Restart independently ...

Illustrative closed-loop episode excerpts for a fish-and-chips task. Three assistants respond to the same wrong-ingredient substitution. Two simulated users continue from the same correction by Assistant B. Frames come from engine-executed actions; dialogue is authored in the style of recorded episodes, not measured output from named models. Recovery excerpts end after discarding food and restarting, not at task completion.

This proof places the figure at the manuscript text width, followed by a 10-point caption. Screenshot details do not replace readable vector labels. The publication PDF preserves text and geometry as vectors; only simulator observations are raster images.