


INTERACTIVE VISUAL GYM / FIGURE PROOF

Environment showcase / wide triptych | 5.5-inch manuscript width

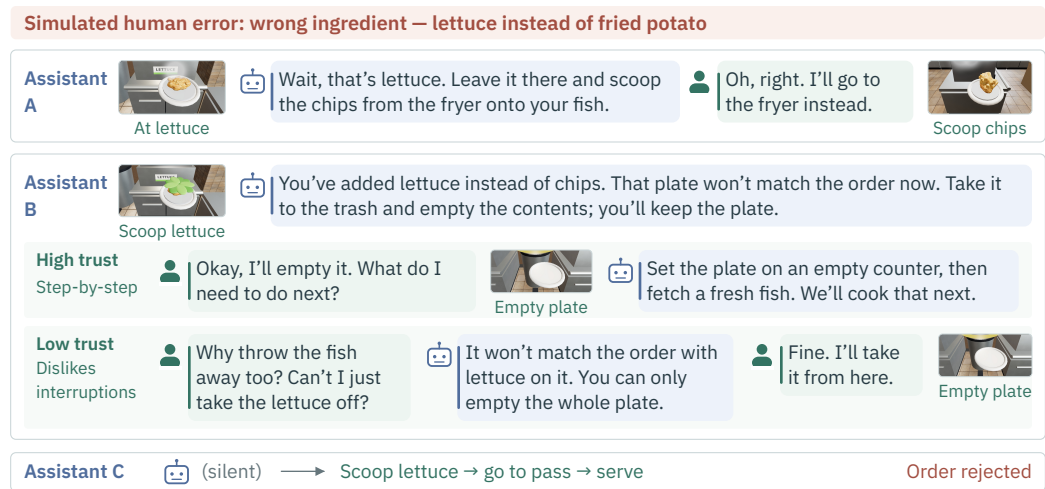


Procedural tasks in three environments. Task descriptions accompany top-down context maps and selected egocentric navigation/action observations. The phone panels retain full screens and show illustrative touch gestures. Ellipses omit intermediate actions; context maps are not execution steps.

This proof places the figure at the manuscript text width, followed by a 10-point caption. Screenshot details do not replace readable vector labels. The publication PDF preserves text and geometry as vectors; only simulator observations are raster images.

INTERACTIVE VISUAL GYM / FIGURE PROOF

Interaction episodes / horizontal strips | 5.5-inch manuscript width

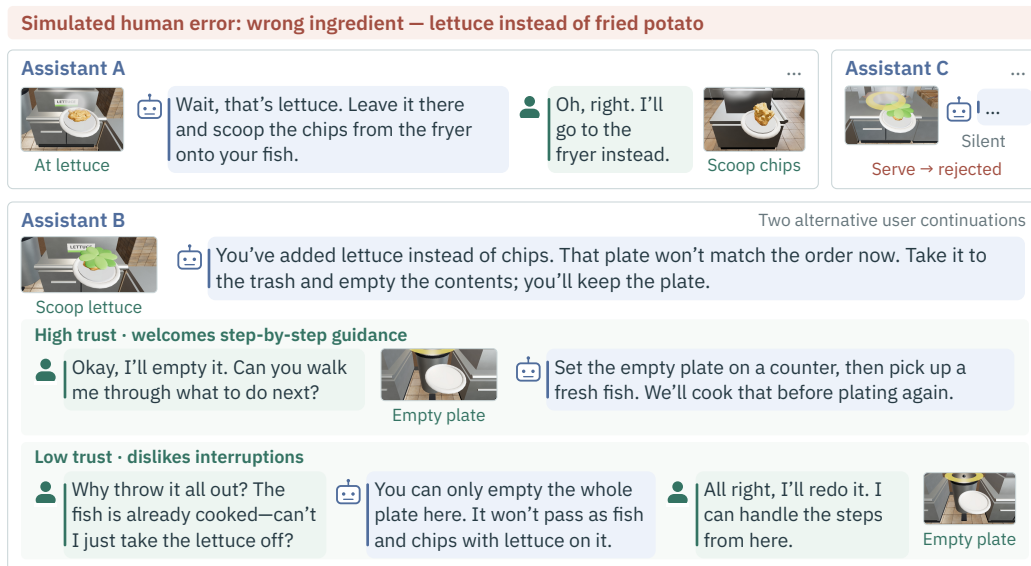


Illustrative episode excerpts for a hard-map fish-and-chips task. Three assistants respond differently to the same wrong-ingredient event. Only Assistant B expands into two alternative user continuations. Speech and user actions are read left to right within each strip. Frames are engine-rendered; dialogue is authored in the style of recorded episodes, not measured model output. The recovery excerpts show emptying the plate, not completed recovery.

This proof places the figure at the manuscript text width, followed by a 10-point caption. Screenshot details do not replace readable vector labels. The publication PDF preserves text and geometry as vectors; only simulator observations are raster images.

INTERACTIVE VISUAL GYM / FIGURE PROOF

Interaction episodes / fuller dialogue alternative | 5.5-inch manuscript width



Illustrative episode excerpts with fuller user replies. Three assistants respond differently to a wrong-ingredient event; Assistant B has two user continuations. Read each continuation from left to right. Dialogue is authored, not a quotation or model-performance result. Only the user acts; recovery observations show the plate emptied, not task completion.

This proof places the figure at the manuscript text width, followed by a 10-point caption. Screenshot details do not replace readable vector labels. The publication PDF preserves text and geometry as vectors; only simulator observations are raster images.