

Assistant	In-time success	Interaction quality					
		Factual grounding	Situational relevance	Actionable guidance	User intent uptake	Guidance conciseness	Overall
Frontier API models							
Gemini 3.8 Flash	30.6	82.7	79.6	68.1	76.2	89.0	53.4
GPT-6 Astra	26.5	86.5	89.4	78.5	97.1	86.1	63.5
Gemini 3.1 Pro	20.8	61.5	66.0	56.6	80.4	65.4	27.9
Real-time streaming API models							
GPT-Realtime 2.1	21.2	80.9	87.7	72.6	76.2	58.1	49.5
Gemini 3.8 Live	11.9	74.7	78.2	58.0	77.7	86.0	39.9
Open-weight non-streaming VLMs							
Muse Glimmer	23.8	59.2	75.3	56.8	65.0	41.3	27.3
Qwen3.8-27B	18.4	70.9	81.7	62.4	85.9	71.9	37.4
Kimi-VL-A3B-Thinking	15.2	34.2	39.4	29.1	39.0	26.0	5.7
InternVideo3-8B	14.5	65.7	45.3	35.0	41.7	68.9	20.9
Molmo2-8B	13.4	54.5	31.7	32.3	33.4	19.2	11.6
Open-weight streaming VLMs							
Qwen3-Omni-30B	26.2	65.8	63.8	51.2	54.3	67.6	25.4
MiniCPM-o 4.5	13.7	55.4	44.7	33.5	46.5	51.8	14.6
Proact-VL	11.3	66.4	25.6	17.2	28.4	44.5	15.8
StreamingVLM	10.9	75.8	25.0	34.8	34.4	22.9	24.5

Table 1 Visual assistance in CookSim. All scores are on a 0–100 scale (↑). In-time success is averaged over three runs. Interaction quality uses run-1 judgments and the full 16-item overall score. Bold/underline indicate the best/second-best assistant in each column.