

	In-time success	Efficiency			Interaction quality					
		Action	Conversation	Overall	Factual grounding	Situational relevance	Actionable guidance	User intent uptake	Guidance conciseness	Overall
Assistant										
Frontier API models										
GPT-6 Astra	90.4	77.2	65.3	71.7	88.4	96.1	93.7	60.0	95.2	75.5
Gemini 3.1 Pro	86.3	80.1	30.0	54.7	91.5	93.8	93.8	33.3	96.2	79.6
Gemini 3.8 Flash	78.5	79.1	59.9	70.3	92.0	96.5	90.5	100.0	98.2	77.8
Real-time streaming API models										
Gemini 3.8 Live	71.9	79.9	21.4	49.2	94.0	91.3	92.2	81.0	96.9	82.6
GPT-Realtime 2.1	66.7	80.4	4.2	41.5	92.1	95.6	90.3	87.0	84.5	76.8
Open-weight non-streaming VLMs										
Qwen3.8-27B	67.8	79.2	12.0	44.9	93.1	94.2	92.9	69.1	96.1	82.0
Muse Glimmer	44.1	75.2	2.9	36.6	80.4	87.5	81.8	65.0	75.2	54.8
Kimi-VL-A3B-Thinking	9.6	75.4	19.7	46.1	71.5	58.5	54.7	37.5	58.9	30.1
Molmo2-8B	0.7	82.6	0.0	36.7	81.5	31.3	38.2	24.0	65.7	15.7
InternVideo3-8B	0.4	63.6	21.4	42.5	70.8	49.3	42.6	27.5	68.5	23.2
Open-weight streaming VLMs										
Qwen3-Omni-30B	15.2	79.8	0.9	35.7	82.2	70.6	72.4	22.0	85.0	49.6
MiniCPM-o 4.5	2.6	81.4	1.2	34.5	75.0	54.8	50.0	49.1	66.4	31.3
Proact-VL	0.0	—	—	—	82.2	10.2	12.5	6.9	43.3	12.4
StreamingVLM	0.0	—	—	—	92.6	14.9	22.4	0.0	30.9	26.9

Table 3 Visual assistance in ScreenSim. All scores are on a 0–100 scale (↑). Success and efficiency are three-run means; efficiency is measured on in-time successes, with persona-weighted overall scores. Rubrics use run 1 and the full 16-item overall score. Bold/underline indicate the best/second-best assistant in each column; dashes denote unavailable scores.