

Assistant	In-time success	Interaction quality					Overall
		Factual grounding	Situational relevance	Actionable guidance	User intent uptake	Guidance conciseness	
Frontier API models							
GPT-6 Astra	90.4	88.4	96.1	93.7	60.0	95.2	75.5
Gemini 3.1 Pro	86.3	91.5	93.8	93.8	33.3	96.2	79.6
Gemini 3.8 Flash	78.5	92.0	96.5	90.5	100.0	98.2	77.8
Real-time streaming API models							
Gemini 3.8 Live	71.9	94.0	91.3	92.2	81.0	96.9	82.6
GPT-Realtime 2.1	66.7	92.1	95.6	90.3	87.0	84.5	76.8
Open-weight non-streaming VLMs							
Qwen3.8-27B	67.8	93.1	94.2	92.9	69.1	96.1	82.0
Muse Glimmer	44.1	80.4	87.5	81.8	65.0	75.2	54.8
Kimi-VL-A3B-Thinking	9.6	71.5	58.5	54.7	37.5	58.9	30.1
Molmo2-8B	0.7	81.5	31.3	38.2	24.0	65.7	15.7
InternVideo3-8B	0.4	70.8	49.3	42.6	27.5	68.5	23.2
Open-weight streaming VLMs							
Qwen3-Omni-30B	15.2	82.2	70.6	72.4	22.0	85.0	49.6
MiniCPM-o 4.5	2.6	75.0	54.8	50.0	49.1	66.4	31.3
Proact-VL	0.0	82.2	10.2	12.5	6.9	43.3	12.4
StreamingVLM	0.0	92.6	14.9	22.4	0.0	30.9	26.9

Table 3 Visual assistance in ScreenSim. All scores are on a 0–100 scale (↑). In-time success is averaged over three runs. Interaction quality uses run-1 judgments and the full 16-item overall score. Bold/underline indicate the best/second-best assistant in each column.