

Assistant	Success	Efficiency			Interaction quality					
	In-time	Action	Conversation	Overall	Factual grounding	Situational relevance	Actionable guidance	User intent uptake	Guidance conciseness	Overall
Frontier API models										
GPT-6 Astra	90.4	77.2	65.3	71.7	88.4	<u>96.1</u>	<u>93.7</u>	60.0	95.2	75.5
Gemini 3.1 Pro	<u>86.3</u>	80.1	30.0	54.7	91.5	93.8	93.8	33.3	96.2	79.6
Gemini 3.8 Flash	78.5	79.1	<u>59.9</u>	<u>70.3</u>	92.0	96.5	90.5	100.0	98.2	77.8
Real-time streaming API models										
Gemini 3.8 Live	71.9	79.9	21.4	49.2	94.0	91.3	92.2	81.0	<u>96.9</u>	82.6
GPT-Realtime 2.1	66.7	80.4	4.2	41.5	92.1	95.6	90.3	<u>87.0</u>	84.5	76.8
Open-weight non-streaming VLMs										
Qwen3.8-27B	67.8	79.2	12.0	44.9	<u>93.1</u>	94.2	92.9	69.1	96.1	<u>82.0</u>
Muse Glimmer	44.1	75.2	2.9	36.6	80.4	87.5	81.8	65.0	75.2	54.8
Kimi-VL-A3B-Thinking	9.6	75.4	19.7	46.1	71.5	58.5	54.7	37.5	58.9	30.1
Molmo2-8B	0.7	82.6	0.0	36.7	81.5	31.3	38.2	24.0	65.7	15.7
InternVideo3-8B	0.4	63.6	21.4	42.5	70.8	49.3	42.6	27.5	68.5	23.2
GLM-4.6V-Flash	0.0	—	—	—	89.2	52.8	56.7	40.9	89.0	42.6
Open-weight streaming VLMs										
Qwen3-Omni-30B	15.2	79.8	0.9	35.7	82.2	70.6	72.4	22.0	85.0	49.6
MiniCPM-o 4.5	2.6	<u>81.4</u>	1.2	34.5	75.0	54.8	50.0	49.1	66.4	31.3
Proact-VL	0.0	—	—	—	82.2	10.2	12.5	6.9	43.3	12.4
StreamingVLM	0.0	—	—	—	92.6	14.9	22.4	0.0	30.9	26.9
Reference conditions										
No assistance	0.0	—	—	—	—	—	—	—	—	—
Scripted reference	89.6	78.4	73.2	76.3	96.2	98.7	94.7	100.0	96.6	85.2

Table 3 Visual assistance in ScreenSim. All scores are on a 0–100 scale (↑). Success and efficiency are three-run means; efficiency is measured on in-time successes, with persona-weighted overall scores. Rubrics use run 1 and the full 16-item overall score. Bold/underline indicate the best/second-best assistant in each column; dashes denote unavailable scores.