

Assistant	Overall quality			Rubric dimensions: three-environment mean				
	CookSim	VHSim	ScreenSim	Factual grounding	Situational relevance	Actionable guidance	User intent uptake	Guidance conciseness
Frontier API models								
Gemini 3.8 Flash	<u>53</u>	46	78	86	86	<u>76</u>	80	<u>89</u>
GPT-6 Astra	64	34	76	84	90	77	<u>79</u>	85
Gemini 3.1 Pro	28	31	80	74	79	71	60	77
Real-time streaming API models								
GPT-Realtime 2.1 (polled)	41	37	77	85	<u>87</u>	74	78	64
Gemini 3.8 Live (polled)	47	42	83	<u>88</u>	81	72	68	90
GPT-Realtime 2.1 (native)	50	—	—	—	—	—	—	—
Gemini 3.8 Live (native)	40	—	—	—	—	—	—	—
Open-weight non-streaming VLMs								
Muse Glimmer	27	26	55	70	83	67	69	54
Qwen3.8-27B	37	29	<u>82</u>	78	85	73	74	76
Kimi-VL-A3B-Thinking	6	12	30	54	53	42	38	39
InternVideo3-8B	21	18	23	67	52	41	31	60
Molmo2-8B	12	17	16	69	40	37	26	42
GLM-4.6V-Flash	47	<u>43</u>	43	90	60	55	50	83
Open-weight streaming VLMs								
Qwen3-Omni-30B	25	30	50	75	73	61	45	77
MiniCPM-o 4.5	15	16	31	64	53	39	45	49
Proact-VL	16	16	12	73	22	17	19	38
StreamingVLM	24	23	27	82	25	29	21	25

Table 2 Interaction quality. Overall quality in each environment and five complementary dimensions of assistance, evaluated with the full CookSim rubric. Scores are on a 0–100 scale (run 1). Overall: full 16-item CookSim rubric with mandatory-failure gating, not the mean of the five displayed dimensions. Dimensions: equal-weight means across all three environments; unavailable for CookSim-only native configurations. Bold / underline: highest / second-highest per column. Row order follows CookSim success within each family.