

Table 3: **Question answering and proactive assistance.** Tracking task history and fulfilling standing requests during cooking episodes.

Assistant	Items	Question answering				Requests
		Temporal order	Action counting	Cooking time	Event recall	Timely alerts
Frontier API models						
GPT-6 Astra	194	50	27	94	91	86
Gemini 3.1 Pro	191	17	13	54	30	46
Gemini 3.8 Flash	201	75	29	60	48	56
Real-time streaming API models						
Gemini 3.8 Live (polled)	191	79	35	59	63	30
GPT-Realtime 2.1 (polled)	202	42	22	54	73	26
Gemini 3.8 Live (native)	154	28	16	41	55	16
GPT-Realtime 2.1 (native)	165	65	24	63	91	10
Open-weight non-streaming VLMs						
Qwen3.8-27B	162	55	5	50	40	27
Molmo2-8B	161	0	0	6	11	18
InternVideo3-8B	140	—	0	23	3	17
GLM-4.6V-Flash	116	—	0	23	14	15
Kimi-VL-A3B-Thinking	115	—	0	26	0	7
Muse Glimmer	115	—	0	41	10	23
Open-weight streaming VLMs						
MiniCPM-o 4.5	108	—	0	14	0	18
Qwen3-Omni-30B	167	17	0	29	11	35
Proact-VL	185	0	0	3	0	3
StreamingVLM	187	0	4	4	0	1

CookSim, run 1; all scores on a 0–100 scale. Questions receive full, partial, or no credit; alerts must fall within the valid response window. Items: scored questions and requests. Dashes: no scored items. Bold: highest score per category.