

Assistant	Items	Question answering				Requests
		Temporal order	Action counting	Cooking time	Event recall	Timely alerts
Frontier API models						
Gemini 3.8 Flash	201	<u>75</u>	<u>29</u>	60	48	<u>56</u>
GPT-6 Astra	194	50	27	94	91	86
Gemini 3.1 Pro	191	17	13	54	30	46
Real-time streaming API models						
GPT-Realtime 2.1 (polled)	202	42	22	54	<u>73</u>	26
Gemini 3.8 Live (polled)	191	79	35	59	63	30
GPT-Realtime 2.1 (native)	165	65	24	<u>63</u>	91	10
Gemini 3.8 Live (native)	154	28	16	41	55	16
Open-weight non-streaming VLMs						
Muse Glimmer	115	—	0	41	10	23
Qwen3.8-27B	162	55	5	50	40	27
Kimi-VL-A3B-Thinking	115	—	0	26	0	7
InternVideo3-8B	140	—	0	23	3	17
Molmo2-8B	161	0	0	6	11	18
GLM-4.6V-Flash	116	—	0	23	14	15
Open-weight streaming VLMs						
Qwen3-Omni-30B	167	17	0	29	11	35
MiniCPM-o 4.5	108	—	0	14	0	18
Proact-VL	185	0	0	3	0	3
StreamingVLM	187	0	4	4	0	1

Table 3 Question answering and proactive assistance. Tracking task history and fulfilling standing requests during cooking episodes. CookSim, run 1; scores on a 0–100 scale. Items: scored questions and requests; dashes: no scored items. Bold / underline: highest / second-highest per category. Rows follow CookSim success within each family. Alerts receive credit only within the valid response window.