

Table S: **Interaction quality in ScreenSim.** Unaggregated scores for layout review.

Assistant	Overall	Factual grounding	Situational relevance	Actionable guidance	User intent uptake	Guidance conciseness
<i>Frontier API models</i>						
GPT-6 Astra	76	88	96	94	60	95
Gemini 3.1 Pro	80	91	94	94	33	96
Gemini 3.8 Flash	78	92	96	91	100	98
<i>Real-time streaming API models</i>						
Gemini 3.8 Live (polled)	83	94	91	92	81	97
GPT-Realtime 2.1 (polled)	77	92	96	90	87	84
<i>Open-weight non-streaming VLMs</i>						
Qwen3.8-27B	82	93	94	93	69	96
Molmo2-8B	16	81	31	38	24	66
InternVideo3-8B	23	71	49	43	28	68
GLM-4.6V-Flash	43	89	53	57	41	89
Kimi-VL-A3B-Thinking	30	71	59	55	37	59
Muse Glimmer	55	80	88	82	65	75
<i>Open-weight streaming VLMs</i>						
MiniCPM-o 4.5	31	75	55	50	49	66
Qwen3-Omni-30B	50	82	71	72	22	85
Proact-VL	12	82	10	12	7	43
StreamingVLM	27	93	15	22	0	31

Scores are on a 0–100 scale (run 1). Overall: full 16-item CookSim rubric with mandatory-failure gating, not the mean of the five displayed dimensions. Dimension scores are averaged within episodes and then across episodes. Bold: highest score per column.