

Assistant	Overall	Factual grounding	Situational relevance	Actionable guidance	User intent uptake	Guidance conciseness
Frontier API models						
Gemini 3.8 Flash	78	92	96	91	100	98
GPT-6 Astra	76	88	96	94	60	95
Gemini 3.1 Pro	80	91	<u>94</u>	94	33	96
Real-time streaming API models						
GPT-Realtime 2.1 (polled)	77	92	96	90	<u>87</u>	84
Gemini 3.8 Live (polled)	83	94	91	92	81	<u>97</u>
Open-weight non-streaming VLMs						
Muse Glimmer	55	80	88	82	65	75
Qwen3.8-27B	<u>82</u>	<u>93</u>	<u>94</u>	<u>93</u>	69	96
Kimi-VL-A3B-Thinking	30	71	59	55	37	59
InternVideo3-8B	23	71	49	43	28	68
Molmo2-8B	16	81	31	38	24	66
GLM-4.6V-Flash	43	89	53	57	41	89
Open-weight streaming VLMs						
Qwen3-Omni-30B	50	82	71	72	22	85
MiniCPM-o 4.5	31	75	55	50	49	66
Proact-VL	12	82	10	12	7	43
StreamingVLM	27	<u>93</u>	15	22	0	31

Table S Interaction quality in ScreenSim. Unaggregated scores for layout review. Scores are on a 0–100 scale (run 1). Overall: full 16-item CookSim rubric with mandatory-failure gating, not the mean of the five displayed dimensions. Dimension scores are averaged within episodes and then across episodes. Bold / underline: highest / second-highest per column. Row order follows CookSim success within each family.